

O(log n) - LOGARITHMIC TIME

The Smart Divider - Divide and Conquer Excellence!

CS 101 - Fall 2025

What is O(log n) - Logarithmic Time?

 The Smart Problem Solver

O(log n) means the algorithm **halves the problem** with each step, creating incredibly efficient performance!

Real-World Analogy:

- Like playing “**20 Questions**” - each question eliminates half the possibilities
- Guessing a number from 1-1000? “Is it > 500?” cuts it in half!
- **32 billion items** → Only **32 steps** to find anything!

Incredible Scaling

- 1,000 items → ~10 steps
- 1,000,000 items → ~20 steps
- 1,000,000,000 items → ~30 steps
- **Mind-blowing efficiency!**

Key Insight The algorithm **eliminates half** the remaining possibilities with each step.

Magic Question: “*Can I eliminate half the data without checking it?*”

If yes → You might achieve O(log n)!

What Makes Algorithms $O(\log n)$?

i The “Divide and Conquer” Strategy

$O(\log n)$ algorithms use smart strategies to avoid checking most of the data.

Binary Search - The Classic

```
def binary_search(sorted_array, target):
    left = 0
    right = len(sorted_array) - 1
    steps = 0

    while left <= right:
        steps += 1
        mid = (left + right) // 2

        if sorted_array[mid] == target:
            return mid, steps # Found it!
        elif sorted_array[mid] < target:
            left = mid + 1      # Search right half
        else:
            right = mid - 1     # Search left half

    return -1, steps # Not found

# Example: Find 7 in [1, 3, 5, 7, 9, 11, 13, 15]
# Step 1: Check middle (7) - Found it!
# Only 1 step for 8 items!
search_space = [i for i in range(16) if i % 2 ==1]
value_to_find = 7
print(f"Search space = {search_space}")
print(f"Value to find = {value_to_find}")
pos, steps = binary_search(search_space,7)
print(f"Position where found = {pos}, steps = {steps}")
# Explain the following
# binary_search([i for i in range(100)],4000)
```

Why $O(\log n)$?

- Each step eliminates half the remaining items
- $\log(n)$ = number of times you can halve n

- Incredibly efficient!

Tree Operations

```
# Binary Search Tree operations
class TreeNode:
    def __init__(self, value):
        self.value = value
        self.left = None
        self.right = None

def tree_search(root, target):
    steps = 0
    current = root

    while current:
        steps += 1
        if current.value == target:
            return True, steps # Found!
        elif target < current.value:
            current = current.left # Go left
        else:
            current = current.right # Go right

    return False, steps # Not found

# Each step eliminates half the tree!
# Perfectly balanced tree: O(log n) guaranteed
```

The Tree Advantage:

- Data is pre-organized for smart searching
- Never need to check more than tree height
- Height $\log(n)$ for balanced trees

The Tree Search In Action

! Important

To use the tree search, a balanced tree first must be created from the data. This may take more time initially, but the search itself is fast. Each additional search will also be time efficient.

```
# Binary Search Tree operations
class TreeNode:
    def __init__(self, value):
        self.value = value
        self.left = None
        self.right = None

    def tree_search(self, target):
        steps = 0
        current = self

        while current:
            steps += 1
            if current.value == target:
                return True, steps # Found!
            elif target < current.value:
                current = current.left # Go left
            else:
                current = current.right # Go right

        return False, steps # Not found

# Each step eliminates half the tree!
# Perfectly balanced tree: O(log n) guaranteed

def insert_node(self, value):
    """Insert a value into the binary search tree"""
    if root is None:
        return TreeNode(value)

    if value < root.value:
        root.left = insert_node(root.left, value)
    else:
        root.right = insert_node(root.right, value)
```

```

    return root

def create_balanced_tree(values):
    """Create a balanced BST from a sorted list"""
    if not values:
        return None

    mid = len(values) // 2
    root = TreeNode(values[mid])

    root.left = create_balanced_tree(values[:mid])
    root.right = create_balanced_tree(values[mid + 1:])

    return root

def print_tree_inorder(root):
    """Print tree values in order (for verification)"""
    if root:
        print_tree_inorder(root.left)
        print(root.value, end=" ")
        print_tree_inorder(root.right)

# Demonstration of O(log n) search
if __name__ == "__main__":
    import math

    print("== Binary Search Tree O(log n) Demonstration ==\n")

    # Create a balanced tree with values 1-15
    values = list(range(1, 16)) # [1, 2, 3, ..., 15]
    root = create_balanced_tree(values)

    print("Created balanced BST with values:", values)
    print("Tree structure (in-order traversal):", end=" ")
    print_tree_inorder(root)
    print("\n")

    # Test searches for different values
    test_values = [1, 5, 8, 12, 15, 20] # Include one that doesn't exist

    print("Search Results:")
    print("-" * 50)

```

```

print(f"{'Value':<6} {'Found':<7} {'Steps':<6} {'Expected Max Steps'}")
print("-" * 50)

n = len(values)
expected_max_steps = math.ceil(math.log2(n)) + 1

for target in test_values:
    found, steps = tree_search(root, target)
    print(f"{target:<6} {'Yes' if found else 'No':<7} {steps:<6} {expected_max_steps}")

print("-" * 50)
print(f"\nFor a balanced BST with {n} nodes:")
print(f"Maximum expected steps: {expected_max_steps} ( log ({n}) + 1 )")
print(f"This demonstrates O(log n) time complexity!")

# Demonstrate with larger tree
print("\n==== Larger Tree Demonstration ====")
large_values = list(range(1, 1001)) # 1000 values
large_root = create_balanced_tree(large_values)

# Test a few searches
large_test_values = [1, 500, 999, 1500] # Include one that doesn't exist
n_large = len(large_values)
expected_max_large = math.ceil(math.log2(n_large)) + 1

print(f"\nTesting with {n_large} nodes:")
print(f"{'Value':<6} {'Found':<7} {'Steps':<6} {'Expected Max Steps'}")
print("-" * 50)

for target in large_test_values:
    found, steps = tree_search(large_root, target)
    print(f"{target:<6} {'Yes' if found else 'No':<7} {steps:<6} {expected_max_large}")

print("-" * 50)
print(f"Maximum expected steps for {n_large} nodes: {expected_max_large} ( log ({n_large}) + 1 )")
print("Notice how the steps remain very small even with 1000 nodes!")

```

Interactive O(log n) Binary Search Demo

 Watch the Halving Magic!

See how binary search eliminates half the possibilities with each step. Try finding different numbers!

Python O(log n) Examples - The Efficient Ones!

Binary Search Implementation

```
import bisect # Python's binary search module

# Manual binary search
def binary_search_manual(arr, target):
    left, right = 0, len(arr) - 1
    steps = 0

    while left <= right:
        steps += 1
        mid = (left + right) // 2

        if arr[mid] == target:
            return mid, steps
        elif arr[mid] < target:
            left = mid + 1
        else:
            right = mid - 1

    return -1, steps

# Using Python's bisect module (optimized)
def binary_search_builtin(arr, target):
    pos = bisect.bisect_left(arr, target)
    if pos < len(arr) and arr[pos] == target:
        return pos
    return -1

# Performance comparison
numbers = list(range(0, 100000, 7)) # 0, 7, 14, 21, ...
target = 9961
```

```

# Manual version
position, steps = binary_search_manual(numbers, target)
print(f"Manual: Found at {position} in {steps} steps")

# Built-in version (also  $O(\log n)$ )
position = binary_search_builtin(numbers, target)
print(f"Built-in: Found at {position}")

```

Heap Operations

```

import heapq

# Min-heap operations - all  $O(\log n)$ 
heap = []

# Insert elements -  $O(\log n)$  each
for value in [64, 34, 25, 12, 22, 11, 90]:
    heapq.heappush(heap, value)      #  $O(\log n)$ 
    print(f"Inserted {value}, heap: {heap}")

# Extract minimum -  $O(\log n)$ 
while heap:
    min_value = heapq.heappop(heap)  #  $O(\log n)$ 
    print(f"Removed {min_value}, remaining: {heap}")

# Why  $O(\log n)$ ?
# Heap is a binary tree structure
# Height =  $\log(n)$ 
# Insert/delete only travel up/down one path
# Path length = tree height =  $O(\log n)$ 

# Priority queue example
class PriorityQueue:
    def __init__(self):
        self.heap = []

    def push(self, item, priority):
        heapq.heappush(self.heap, (priority, item))  #  $O(\log n)$ 

    def pop(self):
        return heapq.heappop(self.heap)[1]  #  $O(\log n)$ 

```

💡 Pro Tip

Python's `bisect` module provides highly optimized $O(\log n)$ operations for sorted lists!

O(log n) vs Other Complexities

Performance Comparison : Champion-like Qualities!

Data Size	$O(1)$	$O(\log n)$	$O(n)$	$O(n^2)$
10	1	3	10	100
100	1	7	100	10,000
1,000	1	10	1,000	1,000,000
1,000,000	1	20	1,000,000	1,000,000,000,000
1,000,000,000	1	30	1,000,000,000	∞ (impractical)

The $O(\log n)$ Sweet Spot:

- Almost as fast as $O(1)$
- Dramatically better than $O(n)$
- Scales beautifully with big data
- Requires organized/sorted data

When $O(\log n)$ Shines

Database Lookups!

```
# Pseudocode

# Database indexing - O(log n) lookups
# Even with millions of records!

# File system searches
# Modern filesystems use B-trees (O(log n))

# Sorted data structures
sorted_students = ["Alice", "Bob", "Charlie", ...] # 10,000 students
import bisect

# O(log n) - incredibly fast!
position = bisect.bisect_left(sorted_students, "Emma")
```

```

# Compare with O(n) linear search
position = sorted_students.index("Emma") # Much slower!

# Geographic searches (quad-trees)
# GPS navigation systems
# Image processing (pyramid algorithms)
# Game development (spatial partitioning)

# The secret: Smart data organization
# pays huge dividends in search speed!

```

Real-World O(log n) Applications

Where You Use O(log n) Every Day!

O(log n) algorithms power many systems you interact with daily.

Tech You Use Daily

```

# Pseudocode

# Database queries with indexes
# When you search your emails, photos, contacts
SELECT * FROM emails WHERE subject LIKE '%meeting%'
# Database uses B-tree index: O(log n)

# Autocomplete systems
def autocomplete(prefix, word_list):
    # Binary search to find start position
    start = bisect.bisect_left(word_list, prefix) # O(log n)
    results = []

    for i in range(start, len(word_list)):
        if word_list[i].startswith(prefix):
            results.append(word_list[i])
        else:
            break
    return results

# Version control systems (Git)
# Git uses binary search to find bugs

```

```
git bisect start      # Start binary search
git bisect bad        # Current version has bug
git bisect good v1.0  # v1.0 was good
# Git automatically finds the problematic commit!
```

Behind-the-Scenes Magic

```
# Pseudocode

# Memory management
# Operating systems use balanced trees
# to track free memory blocks

# Network routing
# Internet routers use prefix trees
# to find optimal paths: O(log n)

# Graphics and gaming
# Collision detection uses spatial trees
# Ray tracing uses BVH trees

# Machine learning
# Decision trees make predictions
# in O(log n) time

# Example: Simple decision tree
def classify_student_performance(hours_studied):
    if hours_studied >= 10:      # O(1) decision
        if hours_studied >= 15:  # O(1) decision
            return "Excellent"  # Total: O(log n) depth
        return "Good"
    else:
        if hours_studied >= 5:
            return "Average"
        return "Needs Improvement"

# Tree height determines performance
# Balanced tree = O(log n) predictions
```

Partner Activity: The Power of Divide and Conquer!

Work with a partner to discover the incredible scaling properties of logarithmic algorithms.

Increase and decrease the search space to experiment. To start, look for *BEGIN YOUR EXPERIMENTS HERE!* in the code.

Experiment 1: Binary Search Scaling Test

```
import time
import bisect
import math

def create_sorted_data(size):
    # Create sorted data for testing
    return list(range(0, size * 2, 2))  # [0, 2, 4, 6, 8, ...]

def binary_search_manual(arr, target):
    left, right = 0, len(arr) - 1
    steps = 0

    while left <= right:
        steps += 1
        mid = (left + right) // 2

        if arr[mid] == target:
            return mid, steps
        elif arr[mid] < target:
            left = mid + 1
        else:
            right = mid - 1

    return -1, steps

# Partner A: Test sizes [1000, 2000, 4000]
# Partner B: Test sizes [8000, 16000, 32000]
def test_logarithmic_scaling(sizes):
    results = []

    for size in sizes:
        data = create_sorted_data(size)
        target = data[-1]  # Search for last element (worst case)
```

```

# Manual binary search with step counting
start = time.time()
pos, steps = binary_search_manual(data, target)
manual_time = time.time() - start

# Built-in binary search
start = time.time()
builtin_pos = bisect.bisect_left(data, target)
builtin_time = time.time() - start

theoretical_steps = math.ceil(math.log2(size))

results.append({
    'size': size,
    'actual_steps': steps,
    'theoretical_steps': theoretical_steps,
    'manual_time': manual_time,
    'builtin_time': builtin_time
})

print(f"Size {size:5d}: {steps:2d} steps (theory: {theoretical_steps:2d}), {manual_t

return results

# -----
# BEGIN YOUR EXPERIMENTS HERE!
# Your partner assignment:
my_sizes = [1000, 2000, 4000] # Change based on assignment
results = test_logarithmic_scaling(my_sizes)

# Partner Discussion:
# 1. How close were actual steps to theoretical log (n)?
# 2. What happened to search time as size doubled?
# 3. Compare with your partner's larger sizes - what pattern emerges?

```

Experiment 2: Heap vs Linear Priority Queue

```

import heapq
import time
import random

```

```

class LinearPriorityQueue:
    """Naive O(n) priority queue for comparison"""
    def __init__(self):
        self.items = []

    def push(self, item, priority):
        self.items.append((priority, item))

    def pop(self):
        if not self.items:
            return None
        # Find minimum priority (O(n))
        min_idx = min(range(len(self.items)), key=lambda i: self.items[i][0])
        return self.items.pop(min_idx)[1]

    def compare_priority_queues(num_operations):
        # Generate random tasks with priorities
        tasks = [(random.randint(1, 100), f"Task_{i}") for i in range(num_operations)]

        # Test heap-based priority queue O(log n)
        heap_pq = []
        start = time.time()

        for priority, task in tasks:
            heapq.heappush(heap_pq, (priority, task))

        processed_heap = []
        while heap_pq:
            processed_heap.append(heapq.heappop(heap_pq)[1])

        heap_time = time.time() - start

        # Test linear priority queue O(n)
        linear_pq = LinearPriorityQueue()
        start = time.time()

        for priority, task in tasks:
            linear_pq.push(task, priority)

        processed_linear = []
        while linear_pq.items:
            processed_linear.append(linear_pq.pop())

```

```

linear_time = time.time() - start

return heap_time, linear_time

# -----
# BEGIN YOUR EXPERIMENTS HERE!
# Partner A: Test with 500 operations
# Partner B: Test with 2000 operations
my_operations = 500 # Change based on assignment

heap_time, linear_time = compare_priority_queues(my_operations)

print(f"Operations: {my_operations}")
print(f"Heap PQ ( $O(\log n)$ ): {heap_time:.4f}s")
print(f"Linear PQ ( $O(n)$ ): {linear_time:.4f}s")
print(f"Heap is {linear_time/heap_time:.1f}x faster!")

# Partner Discussion:
# - How did the performance gap change with more operations?
# - What would happen with 10,000 operations?

```

Group Activity: Logarithmic Thinking Challenge!

Group Problem-Solving: When to Use $O(\log n)$!

Important

Solve these real-world scenarios using logarithmic thinking!

Option 1: Library Book System You are designing a system for a library with 100,000 books. - Books are sorted by ISBN number - Students need to find books quickly - New books are added daily

Questions: 1. How would you implement book lookup? (No code necessary, explain steps) 2. What's the maximum number of steps to find any book? 3. How would this scale to 1,000,000 books?

Option 2: Student Grade Ranking Your school wants to rank 5,000 students by GPA. - GPAs range from 0.0 to 4.0 - Need to quickly find a student's rank - Rankings update when grades change

Questions: 1. How would you store the data for fast ranking lookup? (No code necessary, explain steps) 2. How many steps to find where a 3.7 GPA ranks? 3. What

happens when a student's GPA changes?

Your Turn: Experience $O(\log n)$ Power!

💡 Individual Exploration: Feel the Logarithmic Magic!

After group problem-solving, try these hands-on exercises to experience $O(\log n)$ efficiency.

Exercise 1: Search Race

```
import time
import bisect
import random

# Create test data
size = 100000
sorted_data = sorted([random.randint(1, 1000000) for _ in range(size)])
target = sorted_data[size // 2] # Middle element

# Race 1: Linear search O(n)
start = time.time()
linear_pos = -1
for i, value in enumerate(sorted_data):
    if value == target:
        linear_pos = i
        break
linear_time = time.time() - start

# Race 2: Binary search O(log n)
start = time.time()
binary_pos = bisect.bisect_left(sorted_data, target)
binary_time = time.time() - start + 0.001 # add error

print(f"Linear search (O(n)): {linear_time:.6f}s")
print(f"Binary search (O(log n)): {binary_time:.6f}s")
print(f"Binary search is {linear_time/binary_time:.0f}x faster!")

# Try with different sizes
# Notice how binary search stays fast
# while linear search gets slower
```

Exercise 2: Heap Priority Queue

```
import heapq
import time

# Simulate a hospital emergency room
# Priority queue: lower number = higher priority
emergency_room = []

# Add patients with priorities
patients = [
    (1, "Heart Attack"),      # Highest priority
    (5, "Broken Arm"),
    (2, "Severe Bleeding"),
    (8, "Routine Checkup"),  # Lowest priority
    (3, "Chest Pain"),
    (7, "Headache"),
    (1, "Stroke"),           # Also highest priority
]

print("Adding patients to emergency queue:")
for priority, condition in patients:
    heapq.heappush(emergency_room, (priority, condition))  # O(log n)
    print(f"Added: {condition} (Priority {priority})")

print("\nTreating patients in priority order:")
while emergency_room:
    priority, condition = heapq.heappop(emergency_room)  # O(log n)
    print(f"Treating: {condition} (Priority {priority})")

# Each operation is O(log n)
# Total time: O(n log n) for n operations
# Much better than sorting repeatedly: O(n^2 log n)!
```

If you want to know more about how the Heap Priority Queue code works, check out the Supplemental Slides [10_heap_sorting_slides](#)

The Math Behind O(log n)

i Understanding Logarithms

Don't worry - the math is simpler than you think!

What is $\log(n)$?

```
# log (n) = "How many times can I divide n by 2?"\n\ndef calculate_steps(n):\n    steps = 0\n    while n > 1:\n        n = n // 2 # Divide by 2\n        steps += 1\n    return steps\n\n# Examples:\nprint(f"log (8)  {calculate_steps(8)}")      # 3 steps\nprint(f"log (16)  {calculate_steps(16)}")      # 4 steps\nprint(f"log (1024)  {calculate_steps(1024)}")  # 10 steps\n\n# Binary search does exactly this!\n# Each step eliminates half the data\n# Until only 1 item remains\n\n# Real logarithms:\nimport math\nprint(f"Actual log (1024) = {math.log2(1024)}") # 10.0\nprint(f"Actual log (1000000) = {math.log2(1000000):.1f}")
```

Why Logarithms are Magical

```

linear_steps = n
log_steps = math.ceil(math.log2(n))
speedup = linear_steps / log_steps

print(f"{n:,}\t{linear_steps:,}\t{log_steps}")

# Output shows the dramatic difference!
# Size      O(n) steps    O(log n) steps
# 10          10            4
# 100         100           7
# 1,000       1,000          10
# 10,000      10,000         14
# 100,000     100,000        17
# 1,000,000   1,000,000      20

# Notice: O(log n) barely increases!

```

Achieving $O(\log n)$ - Requirements and Trade-offs

! The $O(\log n)$ Prerequisites

$O(\log n)$ isn't magic - it requires smart data organization!

What You Need for $O(\log n)$

```

# 1. Sorted/Organized Data
unsorted = [64, 34, 25, 12, 22, 11, 90]
# Can't use binary search!

sorted_data = [11, 12, 22, 25, 34, 64, 90]
# Perfect for binary search!

# 2. Tree Structure
# Balanced binary search tree
# Heap (for priority operations)
# B-trees (for databases)

# 3. Divide-and-conquer opportunity
# Can you eliminate half the possibilities?
# If yes,  $O(\log n)$  might be possible

```

```

# Example: Finding square root (Newton's method)
def sqrt_binary_search(n, precision=0.001):
    low, high = 0, n
    while high - low > precision:
        mid = (low + high) / 2
        if mid * mid > n:
            high = mid # Too big, search lower half
        else:
            low = mid # Too small, search upper half
    return (low + high) / 2

```

The Trade-offs

```

# Advantage: Incredible search speed
# Disadvantage: Must maintain organization

# Example: Maintaining a sorted list
sorted_scores = [65, 72, 78, 85, 92]

# Insert new score - O(n) to maintain order!
def insert_sorted(sorted_list, value):
    import bisect
    pos = bisect.bisect_left(sorted_list, value) # O(log n) to find
    sorted_list.insert(pos, value) # O(n) to shift!
    return sorted_list

# Better: Use a balanced tree or heap
import heapq

# For min/max operations, use heap
heap = [65, 72, 78, 85, 92]
heapq.heapify(heap) # O(n) once
heapq.heappush(heap, 80) # O(log n) always!
min_score = heapq.heappop(heap) # O(log n) always!

# Decision guide:
# - Frequent searches, rare updates → Use sorted list
# - Frequent updates → Use heap or balanced tree
# - Need both → Use advanced data structures (B-trees)

```

Advanced O(log n) Patterns

Divide and Conquer Algorithms

```
# Merge sort - O(n log n)
def merge_sort(arr):
    if len(arr) <= 1:
        return arr

    mid = len(arr) // 2
    left = merge_sort(arr[:mid])      # Recursive divide
    right = merge_sort(arr[mid:])    # Recursive divide

    return merge(left, right)         # O(n) merge

# Why O(n log n)?
# - log n levels of recursion (divide by 2 each time)
# - O(n) work at each level
# - Total: O(n) × O(log n) = O(n log n)

# Binary exponentiation - O(log n)
def power_fast(base, exp):
    if exp == 0:
        return 1
    if exp % 2 == 0:
        half = power_fast(base, exp // 2)
        return half * half
    else:
        return base * power_fast(base, exp - 1)

# Calculate 2^100 in only 7 steps instead of 100!
```

Tree Traversal Patterns

```
# Binary search tree operations
class BST:
    def __init__(self, value):
        self.value = value
        self.left = None
        self.right = None

    def insert(self, value):          # O(log n) average
```

```

        if value < self.value:
            if self.left:
                self.left.insert(value)
            else:
                self.left = BST(value)
        else:
            if self.right:
                self.right.insert(value)
            else:
                self.right = BST(value)

    def search(self, value):          # O(log n) average
        if value == self.value:
            return True
        elif value < self.value and self.left:
            return self.left.search(value)
        elif value > self.value and self.right:
            return self.right.search(value)
        return False

    # Segment trees for range queries - O(log n)
    # Fenwick trees for cumulative sums - O(log n)
    # Trie trees for string operations - O(log n) depth

```

Summary: $O(\log n)$ - The Scaling Superstar!

Key Takeaways

$O(\log n)$ - Logarithmic Time is the sweet spot between $O(1)$ and $O(n)$!

What Makes $O(\log n)$ Special - **Incredible scaling** - handles billions of items easily - **Smart elimination** - halves the problem each step - **Requires organization** - data must be structured - **Great trade-off** - almost as fast as $O(1)$, much better than $O(n)$

Python $O(\log n)$ Champions: - Binary search: `bisect` module - Heap operations: `heapq` module - Tree operations: balanced trees - Database queries: indexed lookups

Programming Wisdom

```
# When to choose O(log n):
```

```

# For searching large sorted datasets
import bisect
position = bisect.bisect_left(sorted_data, target) # O(log n)

# For priority queues
import heapq
heapq.heappush(heap, item) # O(log n)
priority_item = heapq.heappop(heap) # O(log n)

# For maintaining sorted order with frequent updates
# Use balanced trees or heaps

# Remember the trade-off:
# - O(1): Instant but needs hash tables
# - O(log n): Nearly instant, needs sorted data
# - O(n): Acceptable for small data, no requirements

# Choose based on your data size and update patterns!

```

Next: Exploring $O(2^n)$ - The Exponential Challenge!

Coming Up Next

- $O(2^n)$ - Exponential Time**
 - * Recursive algorithms and the combinatorial explosion
 - * When algorithms become impractically slow
 - * Interactive Fibonacci and subset generation demos
 - * Dynamic programming to the rescue
 - * Why some problems are inherently difficult
- Questions to Think About:**
 - * When does recursion become dangerous?
 - * How can we optimize exponential algorithms?
 - * What problems are fundamentally hard to solve?
 - * When should we accept “good enough” solutions?

Ready to explore the extreme end of algorithm complexity?